

Import Libraries

```
In [1]: import pandas as pd

In [6]: bostonhousing = pd.read_csv('Boston.csv')

In [8]: bostonhousing

Out[8]:
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
...
501	502	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67	22.4
502	503	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	504	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	505	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.0
505	506	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88	11.9

506 rows × 15 columns

```
Split dataset to X and Y variables

In [10]: Y = bostonhousing.medv
Y

Out[10]:
```

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2
...	...
501	22.4
502	20.6
503	23.9
504	22.0
505	11.9

Name: medv, Length: 506, dtype: float64

```
In [12]: X = bostonhousing.drop(['medv'], axis=1)
X

Out[12]:
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33
...
501	502	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67
502	503	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08
503	504	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64
504	505	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48
505	506	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88

506 rows × 14 columns

```
Data split

In [13]: # Import library

from sklearn.model_selection import train_test_split

Perform 80/20 Data split

In [14]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2)

Data dimension

In [15]: X_train.shape, Y_train.shape

Out[15]: ((404, 14), (404,))

In [16]: X_test.shape, Y_test.shape

Out[16]: ((102, 14), (102,))
```

Linear Regression Model

```
In [18]: # Import library

from sklearn import linear_model
from sklearn.metrics import mean_squared_error, r2_score
```

Build linear regression

```
Defines the regression model

In [19]: model = linear_model.LinearRegression()

Build training model

In [20]: model.fit(X_train, Y_train)

Out[20]: LinearRegression()
```

Apply trained model to make prediction (on test set)

```
In [21]: Y_pred = model.predict(X_test)
```

Prediction results

```
Print model performance

In [22]: print('Coefficients:', model.coef_)
print('Intercept:', model.intercept_)
print('Mean squared error (MSE): %.2f'
      % mean_squared_error(Y_test, Y_pred))
print('Coefficient of determination (R^2): %.2f'
      % r2_score(Y_test, Y_pred))

Coefficients: [-2.63574341e-03 -1.16595019e-01  4.69216810e-02 -2.35131792e-02
 1.86771761e+00 -1.51058363e+01  3.74748184e+00 -5.35028136e-03
-1.58584524e+00  3.33065038e-01 -1.15505207e-02 -9.33238295e-01
 1.04371084e-02 -5.53439477e-01]
Intercept: 36.52638473239881
Mean squared error (MSE): 18.04
Coefficient of determination (R^2): 0.74

String formatting

In [23]: r2_score(Y_test, Y_pred)

Out[23]: 0.7413997618975603

In [24]: r2_score(Y_test, Y_pred).dtype

Out[24]: dtype('float64')

We will be using the modulo operator to format the numbers by rounding it off.

In [25]: '%f' % 0.523810833536016

Out[25]: '0.523811'

In [26]: '%.3f' % 0.523810833536016

Out[26]: '0.524'

We will now round it off to 2 digits

In [27]: '%.2f' % 0.523810833536016

Out[27]: '0.52'
```

Scatter plots

```
In [28]: # Import library

import seaborn as sns

In [29]: # The Data

Y_test

Out[29]:
```

297	20.3
220	26.7
319	21.0
13	20.4
139	17.8
...	...
475	13.3
445	11.8
440	10.5
421	14.2
182	37.9

Name: medv, Length: 102, dtype: float64

```
In [30]: import numpy as np
np.array(Y_test)

Out[30]: array([20.3, 26.7, 21. , 20.4, 17.8, 17. , 15.6, 10.8, 21.7,  8.3, 32. ,
 18. , 46. ,  8.5, 11.3, 23.3, 20.6, 24.8, 13.2, 11. , 26.4, 15.6,
 24.5, 13.1, 24.4, 21. , 24.4, 41.3, 35.2, 24.7, 16.6, 20.2, 24. ,
 24.6, 30.1, 34.9, 13.1, 18.9, 23.6, 17.9, 33. , 25.2, 16.8, 21.1,
  8.8, 32.9, 15.6, 16.1, 31.1, 18.4, 23.2, 22.5,  7.2, 17.3, 50. ,
 20.8, 25. ,  7.5, 20.6, 22. , 22.8, 15.6, 19.5, 18.4, 33.1, 20.1,
 23.2, 28. , 13.8, 19.5, 14.5, 20.5, 14.5, 17.2, 17.6, 50. , 16.2,
 13.6, 21.5, 14.9, 26.2, 21.4, 22.8, 29.1, 14.8, 15.2, 19.4,  8.3,
 19.6, 15. , 24.3, 22. , 14.1, 22.6, 11.8, 21.2, 14.4, 13.3, 11.8,
 10.5, 14.2, 37.9])

In [31]: Y_pred

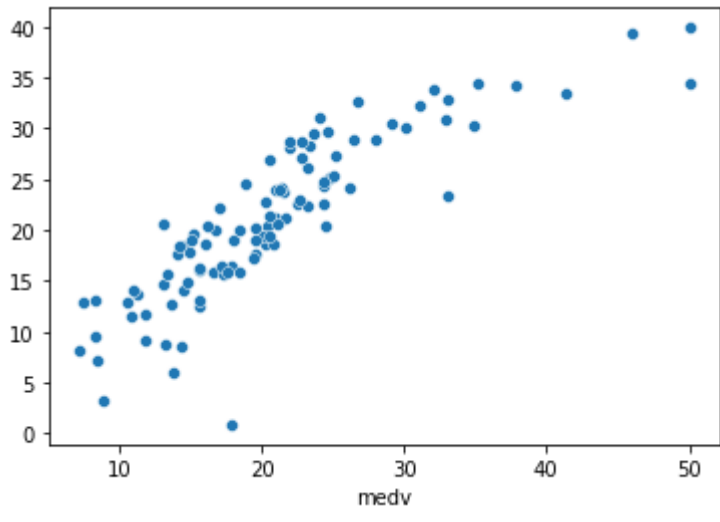
Out[31]: array([18.59935136, 32.59692593, 21.16404792, 20.35480001, 16.45984152,
 22.15793225, 16.05706107, 11.46518899, 21.19663772,  9.45954931,
 33.93529891, 19.02987381, 39.46395322,  7.1791354 , 13.68163892,
 28.39149391, 27.00389682, 25.21334361,  8.77061615, 14.12403152,
 28.87238482, 16.26458202, 20.51219323, 20.62277603, 24.30832501,
 24.01401011, 22.51843377, 33.43313009, 34.53393112, 25.16529841,
 15.85993997, 22.77703804, 31.01301934, 29.66568887, 30.05586673,
 30.36026414, 14.61150457, 24.59581876, 29.53452074,  0.8731239 ,
 23.32875753, 27.39012247, 20.08051769, 20.61785418,  3.24926905,
 30.89568635, 12.58131569, 18.69916412, 32.30355745, 15.82317265,
 26.18577841, 22.5789498 ,  8.07868536, 15.6170231 , 39.02313033,
 18.59786748, 25.262397 , 12.92721025, 21.45973383, 28.0211893 ,
 28.64539803, 13.10583772, 20.1285186 , 19.96503991, 32.87744761,
 19.4022493 , 22.41223786, 28.96790682,  6.06609684, 17.56811559,
 18.46789492, 19.37031407, 14.08799999, 16.41935929, 15.84026005,
 34.4635125 , 20.34331979, 12.7283798 , 23.78371224, 17.89520327,
 24.07329631, 24.256236044, 27.10561578, 30.44579981, 14.96885631,
 19.52889914, 17.22633551, 13.15762302, 19.03355154, 18.95228829,
 24.70417128, 28.78735661, 17.68084657, 22.93320236,  9.14289182,
 23.96986512,  8.54557993, 15.62300708, 11.67537752, 12.83660473,
 18.44855824, 34.24534024])
```

Making the scatter plot

```
In [32]: sns.scatterplot(Y_test, Y_pred)

C:\Users\A21an\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
<AxesSubplot: xlabel='medv'>

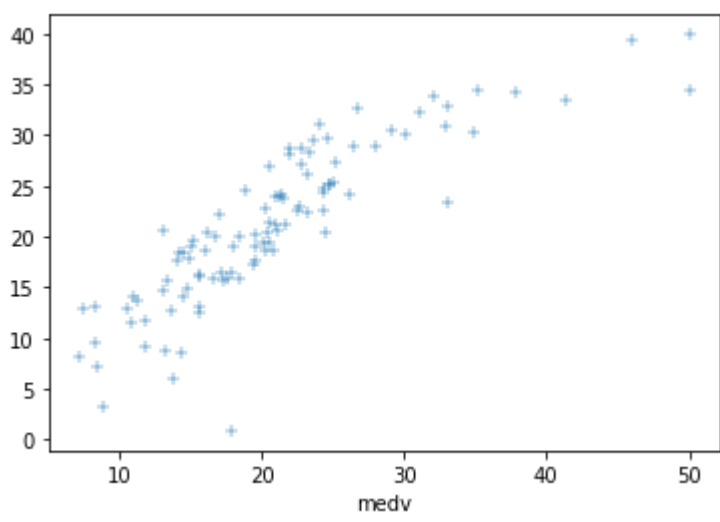
Out[32]:
```



```
In [33]: sns.scatterplot(Y_test, Y_pred, marker="+")

C:\Users\A21an\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
<AxesSubplot: xlabel='medv'>

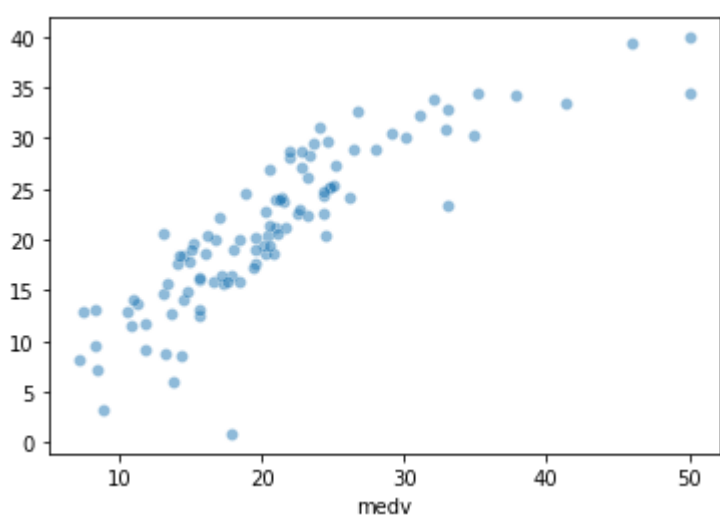
Out[33]:
```



```
In [34]: sns.scatterplot(Y_test, Y_pred, alpha=0.5)

C:\Users\A21an\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
<AxesSubplot: xlabel='medv'>

Out[34]:
```



```
In [ ]:
```